

QUARRY

Autonomous Threat Hunting for the Post-Mythos Era

A governed agentic investigation platform for financial-sector cyber defense.



Confidential

Prepared for strategic partner discussion



APR

07

2026

April 7, 2026 changed the threat-hunting clock.

Frontier agents are now credible enough to compress vulnerability discovery, exploitation, and defensive response cycles.

- Anthropic disclosed Claude Mythos Preview and Project Glasswing, explicitly framing the launch as a security watershed.
- The Federal Reserve later cited Mythos at an FSOC AI / cybersecurity roundtable as evidence of rapidly evolving capability.
- The implication for financial institutions is operational: investigation workflows must move from analyst-hours to machine-minutes while remaining auditable.

The next 6-12 months are the adoption window.

6-12
months

Regulators are turning AI-cyber guidance into operating expectations while attackers gain agentic leverage first.

- Treasury launched public-private AI cybersecurity resources for financial services in 2026.
- Treasury's earlier AI cybersecurity report identified immediate operational risk, cybersecurity, and fraud challenges.
- The buyer will not want a chatbot. The buyer will want governed autonomy: provenance, cost control, reproducibility, and policy guardrails.

Positioning rule: auditable autonomy beats unsupervised autonomy.

Hunters cannot scale linearly with telemetry.

The industry has more alerts, more tools, more data, and fewer people who can connect it all under time pressure.

Telemetry

MORE VOLUME

current SOC reality

Tools

MORE SURFACE

current SOC reality

Hunters

SAME BOTTLENECK

current SOC reality

Reports


LATE EVIDENCE

current SOC reality

- SIEM / EDR / cloud logs are searchable, but the investigation work is still mostly human orchestration.
- Senior hunters are the bottleneck: hypothesis generation, query selection, evidence correlation, and report writing.
- Static detections miss the operating need: turning a vague brief into a defensible investigation trail.

Quarry turns a brief into governed investigation work.

An agentic hunting platform that decomposes the brief, queries telemetry, correlates evidence, and writes the report with a ledger.



Governed

POLICY GUARDRAILS AND READ-ONLY EVALUATION PATHS.
design principle

Traceable

EVERY QUERY AND CONCLUSION GOES TO THE LEDGER.
design principle

Measurable

TIME, TOKENS, COST, AND ACCURACY ARE PRODUCT METRICS.
design principle

The analyst sees a live investigation, not a black box.

Brief, hypotheses, data queries, findings, graph, MITRE coverage, cost, and report generation are visible in one flow



- Input: natural-language investigation brief.
- Process: hypotheses -> connector queries -> evidence ledger -> report.
- Output: analyst-ready narrative with supporting artifacts.

SecRL smoke run: 8/8 scenarios correct.

A pitch-grade regression run against Microsoft's cyber threat investigation benchmark, using the Quarry SQL investigation adapter.

8/8

CORRECT SCENARIOS

one question per incident

8.54s

AVG TIME

brief to answer

95,460

TOKENS

total run

\$0.243

COST

estimated total

- Model: gpt-4o; evaluator: gpt-4o; temperature: 0; max steps: 25.
- Scope note: this is an 8-scenario pitch smoke run, not the full official question-set claim.

Conservative claim: early evidence of fit, not a completed official benchmark certification.

Brazilian fintechs are the first wedge.

Pix created a national-scale real-time payment environment where fraud, identity, mobile, and infrastructure signals converge.

+170M

INDIVIDUAL PIX USERS

BCB public stats

+7B

PIX TRANSACTIONS

January 2026

+R\$3T

MONTHLY VOLUME

October 2025

- Demo persona: FinPlay Pagamentos, a fictitious Sao Paulo fintech with 50,000 active customers.
- Synthetic 30-day dataset covers Pix, mobile banking auth, Open Finance, boleto creation, and injected fraud campaigns.
- Initial use case: Pix fraud and account-takeover investigations that require fast correlation and an auditable report.



Founder-led product, partner-scaled distribution.

The right early team pairs Brazilian market access with global enterprise-security credibility.

Jose

PRODUCT VISION | BRAZILIAN MARKET ACCESS | OPERATOR DISCOVERY

role

Meta partner

ENTERPRISE CREDIBILITY | SECURITY NETWORK LEVERAGE | STRATEGIC DISTRIBUTION

role

Operating model

DEMO FIRST | BENCHMARK VISIBLE | PILOT DISCIPLINED

role

- Jose / Increase Trainer: product vision, Brazilian market access, and hands-on operator discovery.
- Meta co-founder partner: enterprise credibility, security-network leverage, and strategic distribution.
- Operating model: narrow wedge, visible demo, measurable benchmark, then regulated pilot.

BR -> US -> Global, with evidence gates at each step.

The roadmap protects credibility by pairing expansion with benchmark, security, and deployment milestones.



Rule: every expansion step ships with a repeatable demo, benchmark evidence, and operational runbook.

Structure the partnership now.

The near-term ask is not broad fundraising. It is a focused partnership discussion tied to diligence and pilot execution.

- Align on co-founder / advisor / equity structure and operating responsibilities.
- Run 30-day technical diligence: public demo, architecture, SecRL harness, and synthetic fintech dataset.
- Define 60-day pilot path with one Brazilian fintech or MSSP and one enterprise-security design partner.
- Decision point: strategic partnership, seed vehicle, or customer-led buildout.

Week 1

ALIGNMENT

Day 30

DILIGENCE

Day 60

PILOT DECISION

Human-speed teams are entering machine-speed investigations.

Quarry is the control plane that keeps autonomous hunting useful, governed, and defensible.

- Live demo: <https://quarry.12brain.org>
- Backup artifacts: docs/pitch/quarry-v1.pdf and docs/pitch/quarry-v1.pptx
- Next steps: partner diligence, full SecRL question-set run, first regulated pilot.



quarry.12brain.org